

Eigenvalue Decomposition

Predictive Modeling & Statistical Learning

Gaston Sanchez

CC BY-SA 4.0

Matrix Decompositions

Decompositions

Matrix decompositions, also known as matrix factorizations

$$\mathbf{M} = \mathbf{AB} \quad \text{or} \quad \mathbf{M} = \mathbf{ABC}$$

are a means of expressing a matrix as a product of usually two or three **simpler** matrices.

Importance of Decompositions

What for?

Matrix decompositions make it easier to study the properties of matrices. Likewise, many computation tasks become easier with decompositions.

They play a relevant role in multivariate data analysis. Often, the solution to many techniques are obtained (or derived) from a matrix decomposition.

Decompositions: What for?

- ▶ solving systems of linear equations
- ▶ inverting a matrix
- ▶ analyzing numerical stability of a system
- ▶ understanding the structure of data
- ▶ finding basis for column space (or row space) of a matrix

Some Assumptions

Real Matrices

We will assume all matrices to be real matrices, i.e. matrices containing elements in the set of Real numbers.

Dimensions $n \geq p$

Unless otherwise stated, we will also assume matrices with more rows than columns.

Decompositions

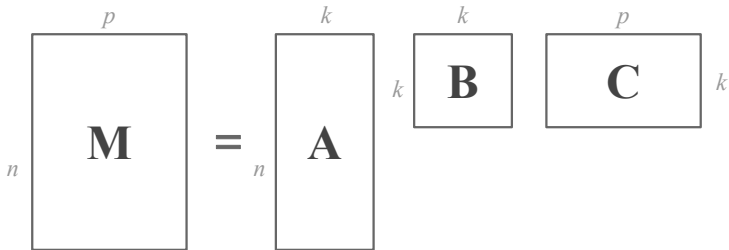
A matrix decomposition can be described by an equation:

$$\mathbf{M} = \mathbf{A}\mathbf{B}\mathbf{C}$$

where the dimensions of the matrices are as follows:

- ▶ \mathbf{M} is $n \times p$ (assume $n > p$)
- ▶ \mathbf{A} is $n \times k$ (usually $k < p$)
- ▶ \mathbf{B} is $k \times k$ (usually diagonal)
- ▶ \mathbf{C} is $k \times p$

Matrix Decomposition



Interpreting Decompositions

The equation that describes a decomposition:

$$\mathbf{M} = \mathbf{ABC}$$

- ▶ does not explain how to compute one
- ▶ does not explain how such decomposition can reveal the structures implicit in a data matrix.
- ▶ Seeing how a matrix decomposition reveals structure in a dataset is more complicated
- ▶ Each decomposition reveals a different kind of implicit structure

Types of matrices

Two types of matrices

We concentrate on the two types of matrices important in statistics:

- ▶ general **rectangular** matrices used to represent data tables.
- ▶ **positive semi-definite** matrices used to represent covariance matrices, correlation matrices, and any matrix that results from a crossproduct.

Two Special Decompositions

EVD and SVD

There are many types of matrix decompositions but for now we are going to consider only two:

- ▶ Eigen-Value Decomposition (EVD)
- ▶ Singular Value Decomposition (SVD)

EVD

Eigenvalue Decomposition

- ▶ EVD applies to square matrices in general.
- ▶ A special type of square matrices are **symmetric** matrices.
- ▶ In data analysis methods, these matrices usually appear in the form of cross-product association matrices:
e.g. $\mathbf{X}^T\mathbf{X}$ and $\mathbf{X}\mathbf{X}^T$
- ▶ The attractive thing about EVD is that when applied to symmetric matrices the results have a “simple” nice structure.

Eigenvalue and Eigenvector

Consider the matrix \mathbf{A} :

$$\mathbf{A} = \begin{bmatrix} 3 & -2 \\ 1 & 0 \end{bmatrix}$$

associated to the linear transformation $T(\mathbf{x})$ given by:

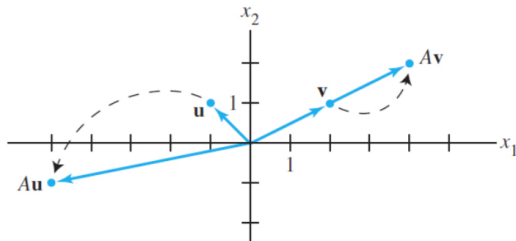
$$T(\mathbf{x}) = \begin{bmatrix} 3 & -2 \\ 1 & 0 \end{bmatrix} \mathbf{x} = \mathbf{A}\mathbf{x}$$

and assume vectors $\mathbf{v} = (2, 1)$ and $\mathbf{u} = (-1, 1)$

Eigenvalue and Eigenvector

$$T(\mathbf{v}) = \begin{bmatrix} 3 & -2 \\ 1 & 0 \end{bmatrix} \begin{bmatrix} 2 \\ 1 \end{bmatrix} = \begin{bmatrix} 4 \\ 2 \end{bmatrix}$$

$$T(\mathbf{u}) = \begin{bmatrix} 3 & -2 \\ 1 & 0 \end{bmatrix} \begin{bmatrix} -1 \\ 1 \end{bmatrix} = \begin{bmatrix} -5 \\ -1 \end{bmatrix}$$



\mathbf{u} is changing its direction, but not \mathbf{v}

Eigenvalue and Eigenvector

Given an $n \times n$ matrix \mathbf{M} , λ is an **eigenvalue** of \mathbf{M} if there exists a non-trivial solution \mathbf{v} of the equation:

$$\mathbf{M}\mathbf{v} = \lambda\mathbf{v}$$

The solution \mathbf{v} is the **eigenvector** associated to the eigenvalue λ

Eigen-Value Decomposition

EVD

An $n \times n$ **symmetric matrix** \mathbf{M} can be decomposed as:

$$\mathbf{M} = \mathbf{A}\mathbf{B}\mathbf{A}^T$$

where

- ▶ \mathbf{A} is a $n \times p$ column **orthonormal** matrix containing the eigen-vectors of \mathbf{M}
- ▶ \mathbf{B} is a $p \times p$ **diagonal** matrix containing the eigen-values of \mathbf{M}

Eigen-Value Decomposition

EVD

A more common notation for EVD is:

$$\mathbf{M} = \mathbf{U}\mathbf{\Lambda}\mathbf{U}^T$$

where

- ▶ \mathbf{U} is a $n \times p$ column **orthonormal** matrix containing the eigen-vectors of \mathbf{M}
- ▶ $\mathbf{\Lambda}$ is a $p \times p$ **diagonal** matrix containing the eigen-values of \mathbf{M}

EVD

$$\mathbf{M} = \mathbf{U}\mathbf{\Lambda}\mathbf{U}^T$$

$$\mathbf{M} = \begin{bmatrix} u_{11} & \cdots & u_{1p} \\ u_{21} & \cdots & u_{2p} \\ \vdots & \ddots & \vdots \\ u_{n1} & \cdots & u_{np} \end{bmatrix} \begin{bmatrix} \lambda_1 & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & \lambda_p \end{bmatrix} \begin{bmatrix} u_{11} & \cdots & u_{n1} \\ u_{12} & \cdots & u_{n2} \\ \vdots & \ddots & \vdots \\ u_{1p} & \cdots & u_{np} \end{bmatrix}$$

Eigenvectors

Vectors, which under a given transformation \mathbf{M} map into themselves or multiples of themselves, are called invariant vectors under that transformation. It follows that such vectors satisfy the relation:

$$\mathbf{M}\mathbf{x} = \lambda\mathbf{x}$$

where λ is a scalar.

Eigenvectors

The matrix equation:

$$\mathbf{M}\mathbf{x} = \lambda\mathbf{x}$$

can be rearranged as follows:

$$\mathbf{M}\mathbf{x} - \lambda\mathbf{x} = \mathbf{0}$$

Eigenvectors

Given

$$\mathbf{M}\mathbf{x} - \lambda\mathbf{x} = \mathbf{0}$$

We can factor out \mathbf{x}

$$(\mathbf{M} - \lambda\mathbf{I})\mathbf{x} = \mathbf{0}$$

Eigenvectors

Obtaining the eigenstructure of a (square) matrix involves solving the **characteristic equation**

$$\det(\mathbf{M} - \lambda_i \mathbf{I}) = 0$$

If \mathbf{M} is of order $n \times n$, then we can obtain n roots of the equation. These roots are called the **eigenvalues**.

EVD in R

`eigen()` in R

`eigen()` function

R provides the function `eigen()` to perform an eigenvalue decomposition of a square matrix.

`eigen()` output

A list with the following components

- ▶ `values` a vector containing the eigenvalues
- ▶ `vectors` a matrix whose columns contain the eigenvectors

EVD example in R

```
# X'X matrix
set.seed(22)
X <- as.matrix(USArrests)
XtX <- t(X) %*% X

# eigenvalue decomposition
EVD = eigen(XtX)

# elements returned by eigen()
names(EVD)

## [1] "values" "vectors"

# vector of eigenvalues
(lambdas = EVD$values)

## [1] 2013735.2431 37957.1103 2084.9578 326.5089
```

EVD example in R (con't)

```
# matrix of eigenvectors
```

```
(V <- EVD$vectors)
```

```
##           [,1]      [,2]      [,3]      [,4]  
## [1,] -0.04239181  0.01616262  0.06588426  0.99679535  
## [2,] -0.94395706  0.32068580 -0.06655170 -0.04094568  
## [3,] -0.30842767 -0.93845891 -0.15496743  0.01234261  
## [4,] -0.10963744 -0.12725666  0.98347101 -0.06760284
```

Properties of Matrix Eigenstructures

Properties of Eigenstructures

1. The sum of the eigenvalues of a matrix \mathbf{A} equals the sum of the main diagonal elements (i.e. the trace) of the matrix.

$$\sum_{i=1}^n \lambda_i = \sum_{i=1}^n a_{ii}$$

2. The product of the eigenvalues of a matrix \mathbf{A} equals the determinant of \mathbf{A}

$$\prod_{i=1}^n \lambda_i = |\mathbf{A}|$$

Properties of Eigenstructures

3. If we have the matrix $\mathbf{B} = \mathbf{A} + k\mathbf{I}$, where k is a scalar, then the eigenvectors of \mathbf{B} are the same as those of \mathbf{A} , and the i -th eigenvalue of \mathbf{B} is

$$\lambda_i + k$$

where λ_i is the i -th eigenvalue of \mathbf{A}

4. If we have the matrix $\mathbf{C} = k\mathbf{A}$, where k is a scalar, then \mathbf{C} has the same eigenvectors as \mathbf{A} and

$$k\lambda_i$$

is the eigenvalue of \mathbf{C} , where λ_i is the i -th eigenvalue of \mathbf{A}

Properties of Eigenstructures

5. If we have the matrix \mathbf{A}^p , where p is a positive integer, then scalar, then \mathbf{A}^p has the same eigenvectors as \mathbf{A} and

$$\lambda_i^p$$

is the i -th eigenvalue of \mathbf{A}^p , where λ_i is the i -th eigenvalue of \mathbf{A}

6. If \mathbf{A}^{-1} exists, then \mathbf{A}^{-p} has the same eigenvectors as \mathbf{A} and

$$\lambda_i^{-p}$$

is the i -th eigenvalue of \mathbf{A}^{-p} corresponding to the i -th eigenvalue of \mathbf{A}

Properties of Eigenstructures

7. If a symmetric matrix \mathbf{A} can be written as the product

$$\mathbf{A} = \mathbf{U}\mathbf{D}\mathbf{U}^T$$

where \mathbf{D} is a diagonal with all entries nonnegative and \mathbf{U} is an orthogonal matrix of eigenvectors, then

$$\mathbf{A}^{1/2} = \mathbf{U}\mathbf{D}^{1/2}\mathbf{U}^T$$

and it is the case that $\mathbf{A}^{1/2}\mathbf{A}^{1/2} = \mathbf{A}$

Properties of Eigenstructures

8. If a symmetric matrix \mathbf{A}^{-1} can be written as the product

$$\mathbf{A}^{-1} = \mathbf{U}\mathbf{D}^{-1}\mathbf{U}^T$$

where \mathbf{D}^{-1} is a diagonal with all entries nonnegative and \mathbf{U} is an orthogonal matrix of eigenvectors, then

$$\mathbf{A}^{-1/2} = \mathbf{U}\mathbf{D}^{-1/2}\mathbf{U}^T$$

and it is the case that $\mathbf{A}^{-1/2}\mathbf{A}^{-1/2} = \mathbf{A}^{-1}$

Power Method

About the Power Method

One of the basic procedures following a successive approximation approach is precisely the **Power Method**.

In its simplest form, the Power Method (PM) allows us to find **the largest** eigenvector and its corresponding eigenvalue.

About the Power Method

Choose an arbitrary vector \mathbf{w}_0 to which we will apply the symmetric matrix \mathbf{S} repeatedly to form the following sequence:

$$\mathbf{w}_1 = \mathbf{S}\mathbf{w}_0$$

$$\mathbf{w}_2 = \mathbf{S}\mathbf{w}_1 = \mathbf{S}^2\mathbf{w}_0$$

$$\mathbf{w}_3 = \mathbf{S}\mathbf{w}_2 = \mathbf{S}^3\mathbf{w}_0$$

$$\vdots$$

$$\mathbf{w}_k = \mathbf{S}\mathbf{w}_{k-1} = \mathbf{S}^k\mathbf{w}_0$$


Power Method: Example

Consider a matrix \mathbf{S}

$$\mathbf{S} = \begin{bmatrix} 3 & 1 \\ 1 & 2 \end{bmatrix}$$


and an initial vector \mathbf{w}_0

$$\mathbf{w}_0 = \begin{bmatrix} -0.4 \\ 1.3 \end{bmatrix}$$



About the Power Method

- ▶ In practice, we must rescale the obtained vector \mathbf{w}_k at each step.
- ▶ The rescaling will allows us to judge whether the sequence is converging.
- ▶ After some iterations, the vector \mathbf{w}_{k-1} and \mathbf{w}_k will be very similar
- ▶ Assuming a reasonable scaling strategy, the sequence will usually converge to the dominant eigenvector of \mathbf{S} .



Dominant Eigenvalue

The obtained vector is the dominant eigenvector. To get the corresponding eigenvalue we calculate the so-called **Rayleigh quotient** given by:

$$\lambda = \frac{\mathbf{w}_k^T \mathbf{S} \mathbf{w}_k}{\mathbf{w}_k^T \mathbf{w}_k}$$

Remarks

Conditions for the power method to be successfully used:

- ▶ The matrix must have a *dominant* eigenvalue.
- ▶ The starting vector \mathbf{w}_0 must be nonzero.
- ▶ We need to scale each of the vectors \mathbf{w}_k otherwise the algorithm will “explode”

PM Pseudocode

Let's consider a more detailed version of the PM algorithm:

1. Start with an arbitrary initial vector \mathbf{w}
2. Obtain product $\tilde{\mathbf{w}} = \mathbf{S}\mathbf{w}$
3. Normalize $\tilde{\mathbf{w}}$

$$\text{e.g. } \mathbf{w} = \frac{\tilde{\mathbf{w}}}{\|\tilde{\mathbf{w}}\|_{p=2}}$$

4. Compare \mathbf{w} with its previous version
5. Repeat steps 2 till 4 until convergence

Why does the PM work?

Assume that the matrix \mathbf{S} has p eigenvalues $\lambda_1, \lambda_2, \dots, \lambda_p$, and that they are ordered in decreasing way

$$|\lambda_1| > |\lambda_2| \geq \dots \geq |\lambda_p|.$$

Note that the first eigenvalue is strictly greater than the second one. This is a very important assumption.

In the same way, we'll assume that the matrix \mathbf{S} has p linearly independent vectors $\mathbf{u}_1, \dots, \mathbf{u}_p$ ordered in such a way that \mathbf{u}_j corresponds to λ_j .

Why does the PM work?

The initial vector \mathbf{w}_0 may be expressed as a linear combination of $\mathbf{u}_1, \dots, \mathbf{u}_p$

$$\mathbf{w}_0 = a_1 \mathbf{u}_1 + \dots + a_p \mathbf{u}_p$$

At every step of the iterative process the vector \mathbf{w}_k is given by:

$$\mathbf{w}_k = a_1 \lambda_1^k \mathbf{u}_1 + \dots + a_p \lambda_p^k \mathbf{u}_p$$

Why does the PM work?

Since λ_1 is the dominant eigenvalue, the component in the direction of \mathbf{u}_1 becomes relatively greater than the other components as k increases. If we knew λ_1 in advance, we could rescale at each step by dividing by it to get:

$$\left(\frac{1}{\lambda_1^k}\right) \mathbf{w}_k = a_1 \mathbf{u}_1 + \cdots + a_p \left(\frac{\lambda_p^k}{\lambda_1^k}\right) \mathbf{u}_p$$

which converges to the eigenvector $a_1 \mathbf{u}_1$, provided that a_1 is nonzero.

Why does the PM work?

Of course, in real life this scaling strategy is not possible—we don't know λ_1 . Consequently, the eigenvector is determined only up to a constant multiple, which is not a concern since the really important thing is the *direction* not the length of the vector.

The speed of the convergence depends on how bigger λ_1 is respect with to λ_2 , and on the choice of the initial vector \mathbf{w}_0 . If λ_1 is not much larger than λ_2 , then the convergence will be slow.

More Remarks

- ▶ The power method is a sequential method.
- ▶ We can obtain $\mathbf{w}_1, \mathbf{w}_2$, and so on, step by step.
- ▶ If we only need the first k vectors, we can stop the procedure at the desired stage.

Obtaining more eigenvectors?

Once we've obtained the first eigenvector \mathbf{w}_1 and eigenvalue λ_1 , we can compute the second eigenvector by reducing the matrix \mathbf{S} by the amount explained by the first eigenvector.

This operation of reduction is called **deflation** and the residual matrix is obtained as:

$$\mathbf{S}_1 = \mathbf{S} - \lambda_1 \mathbf{w}_1 \mathbf{w}_1^T$$

To get the second eigenvalue and its corresponding eigenvector, we operate on \mathbf{S}_1 in the same way as the operations on \mathbf{S} .

References

- ▶ **Multivariate Analysis** by Maurice Tatsuoka (1988). *Chapter 5: More Matrix Algebra*. Macmillan Publishing.
- ▶ **Mathematical Tools for Applied Multivariate Analysis** by J.D. Carroll, P.E. Green, and A. Chaturvedi (1997). *Chapter 5: Decomposition of Matrix Transformations: Eigenstructures and Quadratic Forms*. Academic Press.
- ▶ **Hand-on Matrix Algebra using R** by Hrishikesh Vinod (2011). *Chapter 9: Eigenvalues and Eigenvectors*. World Scientific.